

Locks and Condition Variables

Lecture Notes for CS 140

Spring 2020

John Ousterhout

- Readings for this topic from *Operating Systems: Principles and Practice*: Sections 5.2-5.4.
- Needed: higher-level synchronization mechanism that provides
 - Mutual exclusion: easy to create critical sections
 - Blocking: delay a thread until some desired event occurs

Locks

- **Lock: an object that can only be owned by a single thread at any given time.** Basic operations on a lock:
 - `acquire`: mark the lock as owned by the current thread; if some other thread already owns the lock then first wait until the lock is free. Lock typically includes a queue to keep track of multiple waiting threads.
 - `release`: mark the lock as free (it must currently be owned by the calling thread).
- Too much milk solution with locks (using Pintos APIs):

```
struct lock l;  
...  
lock_acquire(&l);  
if (milk == 0) {  
    buy_milk();  
}  
lock_release(&l);
```
- A more complex example: producer/consumer.
 - Producers add characters to a buffer
 - Consumers remove characters from the buffer
 - Characters will be removed in the same order added
 - Version 1:

```
char buffer[SIZE];
int count = 0, putIndex = 0, getIndex = 0;
struct lock l;
lock_init(&l);

void put(char c) {
    lock_acquire(&l);
    count++;
    buffer[putIndex] = c;
    putIndex++;
    if (putIndex == SIZE) {
        putIndex = 0;
    }
    lock_release(&l);
}

char get() {
    char c;
    lock_acquire(&l);
    count--;
    c = buffer[getIndex];
    getIndex++;
    if (getIndex == SIZE) {
        getIndex = 0;
    }
    lock_release(&l);
    return c;
}
```

- Version 2 (handle empty/full cases):

```

char buffer[SIZE];
int count = 0, putIndex = 0, getIndex = 0;
struct lock l;
lock_init(&l);

void put(char c) {
    lock_acquire(&l);
    while (count == SIZE) {
        lock_release(&l);
        lock_acquire(&l);
    }
    count++;
    buffer[putIndex] = c;
    putIndex++;
    if (putIndex == SIZE) {
        putIndex = 0;
    }
    lock_release(&l);
}

char get() {
    char c;
    lock_acquire(&l);
    while (count == 0) {
        lock_release(&l);
        lock_acquire(&l);
    }
    count--;
    c = buffer[getIndex];
    getIndex++;
    if (getIndex == SIZE) {
        getIndex = 0;
    }
    lock_release(&l);
    return c;
}

```

Condition Variables

- Synchronization mechanisms need more than just mutual exclusion; also need a way to wait for another thread to do something (e.g., wait for a character to be added to the buffer)
- *Condition variables*: used to wait for a particular state to be reached (e.g. characters in buffer).
 - `wait(condition, lock)`: atomically release lock, put thread to sleep until condition is signaled; when thread wakes up again, re-acquire lock before returning.
 - `signal(condition, lock)`: if any threads are waiting on condition, wake up one of them. Caller must hold lock, which must be the same as the lock used in the wait call.
 - `broadcast(condition, lock)`: same as signal, except wake up all waiting threads.
 - Note: after signal, signaling thread keeps lock, waking thread goes on the queue waiting for the lock.

- Warning: when a thread wakes up after `cond_wait` there is no guarantee that the desired condition still exists: another thread might have snuck in.

- Producer/Consumer, version 3 (with condition variables):

```
char buffer[SIZE];
int count = 0, putIndex = 0, getIndex = 0;
struct lock l;
struct condition charAdded;
struct condition charRemoved;

lock_init(&l);
cond_init(&charAdded);
cond_init(&charRemoved);

void put(char c) {
    lock_acquire(&l);
    while (count == SIZE) {
        cond_wait(&charRemoved, &l);
    }
    count++;
    buffer[putIndex] = c;
    putIndex++;
    if (putIndex == SIZE) {
        putIndex = 0;
    }
    cond_signal(&charAdded, &l);
    lock_release(&l);
}

char get() {
    char c;
    lock_acquire(&l);
    while (count == 0) {
        cond_wait(&charAdded, &l);
    }
    count--;
    c = buffer[getIndex];
    getIndex++;
    if (getIndex == SIZE) {
        getIndex = 0;
    }
    cond_signal(&charRemoved, &l);
    lock_release(&l);
    return c;
}
```

Monitors

- When locks and condition variables are used together like this, the result is called a *monitor*:
 - A shared data structure
 - A collection of procedures
 - One lock that must be held whenever accessing the shared data (typically each procedure acquires the lock at the very beginning and releases the lock before returning).

- One or more condition variables used for waiting.
- There are other synchronization mechanisms besides locks and condition variables. Be sure to read about semaphores in the book or in the Pintos documentation.